

CS330: Operating Systems

Semaphore, Classical problems

Semaphores

- Mutual exclusion techniques allows exactly one thread to access the critical section which can be restrictive
- Consider a scenario when a finite array of size N is accessed from a set of producer and consumer threads. In this case,
 - At most N concurrent producers are allowed if array is empty
 - At most N concurrent consumers are allowed if array is full
 - If we use mutual exclusion techniques, only one producer or consumer is allowed at any point of time

Operations on semaphore

```
struct semaphore{  
    int value;  
    spinlock_t *lock;  
    queue *waitQ;  
}sem_t;
```

// Operations

```
sem_init(sem_t *sem, int init_value);  
sem_wait(sem_t *sem);  
sem_post(sem_t *sem);
```

- Semaphores can be initialized by passing an initial value
- *sem_wait* waits (if required) till the value becomes +ve and returns after decrementing the value
- *sem_post* increments the value and wakes up a waiting context
- Other notations: P-V, down-up, wait-signal

Unix semaphores

```
#include <semaphore.h>
```

```
main(){  
    sem_t s;  
    int K = 5;  
    sem_init(&s, 0, K);  
    sem_wait(&s);  
    sem_post(&s);  
}
```

- Can be used to in a multi-threaded process or across multiple processes
- If second argument is 0, the semaphore can be used from multiple threads
- Semaphores initialized with value = 1 (third argument) is called a binary semaphore and can be used to implement *blocking(waiting) locks*
- Initialize: `sem_init(s, 0, 1)`
lock: `sem_wait(s)`, unlock: `sem_post(s)`

Semaphore usage example: wait for child

```
child(){  
    ...  
    sem_post(s);  
    exit(0);  
}  
int main (void ){  
    sem_init(s, 0);  
    if(fork() == 0)  
        child();  
    sem_wait(s);  
}
```

- Assume that the semaphore is accessible from multiple processes, value initialized to zero
- If parent is scheduled after the child creation, it waits till child finishes
- If child is scheduled and exits before parent, parent does not wait for the semaphore

Semaphore usage example: ordering

A=0; B=0;

Thread-0 {

 A = 1;

 printf("B = %d\n", B);

}

- What are the possible outputs?
- (A = 1, B= 1), (A = 1, B = 0), (A = 0, B=1)
- How to guarantee A = 1, B= 1?

Thread-1 {

 B=1;

 printf("A = %d\n", A);

}

Semaphore usage example: ordering

```
sem_init(&s1, 0);  
A=0; B=0;  
Thread - 0 {  
    A = 1;  
    sem_wait(&s1);  
    printf("B = %d\n", B);  
}  
Thread - 1 {  
    B=1;  
    sem_post(&s1);  
    printf("A = %d\n", A);  
}
```

- What are the possible outputs?
- (A = 1, B = 1), (A=0, B=1)
- How to guarantee A = 1, B= 1?

Ordering with two semaphores

```
sem_init(s1, 0);  
sem_init(s2, 0);  
A=0; B=0;
```

- Waiting for each other guarantees
desired output

Thread - 0

```
{  
    A = 1;  
    sem_post(s1);  
    sem_wait(s2);  
    printf("%d\n", B);  
}
```

Thread - 1

```
{  
    B=1;  
    sem_wait(s1);  
    sem_post(s2);  
    printf("%d\n", A);  
}
```


Producer-consumer problem

```
DoProducerWork(){
```

```
while(1){
```

```
    item_t item = prod_p();
```

```
    produce(item);
```

```
}
```

```
}
```



```
DoConsumerWork(){
```

```
while(1){
```

```
    item_t item = consume();
```

```
    cons_p(item);
```

```
}
```

```
}
```

- A buffer of size N, one or more producers and consumers
- Producer produces an element into the buffer (after processing)
- Consumer extracts an element from the buffer and processes it
- Example: A multithreaded web server, network protocol layers etc.
- How to solve this problem using semaphores?

Buggy #1

```
item_t A[n], pctr=0, cctr = 0;  
sem_t empty = sem_init(n), used = sem_init(0);
```

```
produce(item_t item){  
    sem_wait(&empty);  
    A[pctr] = item;  
    pctr = (pctr + 1) % n;  
    sem_post(&used);  
}
```

```
item_t consume() {  
    sem_wait(&used);  
    item_t item = A[cctr];  
    cctr = (cctr + 1) % n;  
    sem_post(&empty);  
    return item;  
}
```

- This solution does not work. What is the issue?
- The counters (pctr and cctr) are not protected, can cause race conditions

Buggy #2

```
item_t A[n], pctr=0, cctr = 0; lock_t *L = init_lock();  
sem_t empty = sem_init(n), used = sem_init(0);
```

```
produce(item_t item){  
    lock(L); sem_wait(&empty);  
    A[pctr] = item;  
    pctr = (pctr + 1) % n;  
    sem_post(&used); unlock(L);  
}
```

```
item_t consume( ) {  
    lock(L); sem_wait(&used);  
    item_t item = A[cctr];  
    cctr = (cctr + 1) % n;  
    sem_post(&empty); unlock(L);  
    return item;  
}
```

- What is the problem?
- Consider empty = 0 and producer has taken lock before the consumer. This results in a deadlock, consumer waits for L and producer for empty

A working solution

```
item_t A[n], pctr=0, cctr = 0; lock_t *L = init_lock();  
sem_t empty = sem_init(n), used = sem_init(0);
```

```
produce(item_t item){  
    sem_wait(&empty); lock(L);  
    A[pctr] = item;  
    pctr = (pctr + 1) % n;  
    unlock(L); sem_post(&used);  
}
```

```
item_t consume() {  
    sem_wait(&used); lock(L)  
    item_t item = A[cctr];  
    cctr = (cctr + 1) % n;  
    unlock(L); sem_post(&empty);  
    return item;  
}
```

- The solution is deadlock free and ensures correct synchronization, but very much serialized (inside produce and consume)
- What if we use separate locks for producer and consumer?

Solution with separate mutexes

```
item_t A[n], pctr=0, cctr = 0; lock_t *P = init_lock(), *C=init_lock();  
sem_t empty = sem_init(n), used = sem_init(0);
```

```
produce(item_t item){  
    sem_wait(&empty); lock(P);  
    A[pctr] = item;  
    pctr = (pctr + 1) % n;  
    unlock(P); sem_post(&used);  
}
```

```
item_t consume() {  
    sem_wait(&used); lock(C)  
    item_t item = A[cctr];  
    cctr = (cctr + 1) % n;  
    unlock(C); sem_post(&empty);  
    return item;  
}
```

- Does this solution work?
- Homework: Assume that item is a large object and copy of item takes long time. How can we perform the copy operation without holding the lock?